

MysecondProject

JOSE IGNACIO SORDO SIERPE

7/1/2022

0.1 1. Introduction

As the second project of the Edx Data Science Capstone course called "Choose Your Own", I am going to choose the data set called PimaIndiansDiabetes

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases, and is available at <https://www.kaggle.com/uciml/pima-indians-diabetes-database>

This dataset tries to predict the probability that a patient has diabetes or not based on certain diagnostic measures included in the dataset.

There are several restrictions in the database being taken in this case, that all patients have to be female, over 21 years of age and of Pima Indian descent.

In the first place we will make a statistical and graphic description of the data, I will analyze if there are statistically significant differences between the two groups (positive and negative) in each of the predictor variables.

Supervised learning machine learning techniques will be applied using 7 different binary classification algorithms and in the results phase we will stick with those that achieve better accuracy.

0.2 2. Methodology and General Description

First we will load the data set and the necessary libraries for the project.

```
# load libraries
library(mlbench)
```

```
## Warning: package 'mlbench' was built under R version 4.0.5
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.0.5
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 4.0.5
```

```
## Loading required package: lattice
```

```
library(Hmisc)
```

```
## Warning: package 'Hmisc' was built under R version 4.0.5
```

```
## Loading required package: survival
```

```
## Warning: package 'survival' was built under R version 4.0.5
```

```
##
## Attaching package: 'survival'
```

```
## The following object is masked from 'package:caret':
##
##   cluster
```

```
## Loading required package: Formula
```

```
##
## Attaching package: 'Hmisc'
```

```
## The following objects are masked from 'package:base':  
##  
##   format.pval, units
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 4.0.5
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
##   margin
```

```
library(Matrix)  
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.0.5
```

```
## Loaded glmnet 4.1-3
```

```
library (pROC)
```

```
## Warning: package 'pROC' was built under R version 4.0.5
```

```
## Type 'citation("pROC")' for a citation.
```

```
##  
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':  
##  
##   cov, smooth, var
```

```
library(ROCR)
```

```
## Warning: package 'ROCR' was built under R version 4.0.5
```

```
library(funModeling)
```

```
## Warning: package 'funModeling' was built under R version 4.0.5
```

```
## funModeling v.1.9.4 :)  
## Examples and tutorials at livebook.datascienceheroes.com  
## / Now in Spanish: librovivodecienciadedatos.ai
```

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.0.5
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:randomForest':  
##  
##   combine
```

```
## The following objects are masked from 'package:Hmisc':  
##  
## src, summarize
```

```
## The following objects are masked from 'package:stats':  
##  
## filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
library(smotefamily)
```

```
## Warning: package 'smotefamily' was built under R version 4.0.5
```

```
library(mlbench)  
library(ggpubr)
```

```
## Warning: package 'ggpubr' was built under R version 4.0.5
```

```
library(onewaytests)
```

```
## Warning: package 'onewaytests' was built under R version 4.0.5
```

```
##  
## Attaching package: 'onewaytests'
```

```
## The following object is masked from 'package:Hmisc':  
##  
## describe
```

```
# load data  
data(PimaIndiansDiabetes)  
# rename dataset to keep code below generic  
dataset <- PimaIndiansDiabetes
```

The methodology will follow the following successive steps: 1.- Description of the data set. 2.- Obtaining the statistical descriptions. 3.- Numerical and graphic analysis of the variables to check if they follow a normal distribution. 4.- Analysis of differences in means and variances to check if there are statistically significant differences between the variables. 5.- Generate a partition of 70% of the sample for training, reserving 30% for validation. 6.- A cross validation will be applied to the 70% partition. 7.- Training of the different algorithms.

0.2.1 2.1. Dataset Description

```
#dataset description.
```

```
head(dataset)
```

```
## pregnant glucose pressure triceps insulin mass pedigree age diabetes  
## 1      6   148    72   35    0 33.6  0.627  50    pos  
## 2      1    85    66   29    0 26.6  0.351  31    neg  
## 3      8   183    64    0    0 23.3  0.672  32    pos  
## 4      1    89    66   23   94 28.1  0.167  21    neg  
## 5      0   137    40   35  168 43.1  2.288  33    pos  
## 6      5   116    74    0    0 25.6  0.201  30    neg
```

The data set is composed of 768 observations with 9 different variables. The variables are: Pregnant — Number of times pregnant Glucose — Plasma glucose concentration a 2 hours in an oral glucose tolerance test Pressure — Diastolic blood pressure (mm Hg) Triceps — Triceps skin fold thickness (mm) Insulin — 2-Hour serum insulin (mu U/ml) Mass — Body mass index (weight in kg/(height in m)^2) Pedigree — Diabetes pedigree function Age — Age (years) Diabetes — Class variable (positive or negative)

0.2.2 2.2. Descriptive Statistical

```
# cuantitative variables
df = dataset[,c(1:8)]

# statistical descriptives
estadisticaDescriptiva <- t(do.call(data.frame,
list(mean = apply(df, 2, mean),
Desv.Estandar = apply(df, 2, sd),
Mediana = apply(df, 2, median),
IQR = apply(df,2,IQR),
Min = apply(df, 2, min),
Max = apply(df, 2, max),
Rango = apply(df, 2,max)-apply(df, 2,min),
Cuartil1 = apply(df,2,quantile,probs = c(0.25)),
Cuartil3 = apply(df,2,quantile,probs = c(0.75)),
N = apply(df,2,length),
ErrorEstandar = apply(df,2,sd)/sqrt(apply(df,2,length)),
IC95MediaLower = apply(df,2,mean)-1.96*apply(df,2,sd)/sqrt(apply(df,2,length)),
IC95MediaUpper = apply(df,2,mean)+1.96*apply(df,2,sd)/sqrt(apply(df,2,length)),
Varianza = apply(df, 2, var)
)))

estadisticaDescriptiva
```

```
##      pregnant  glucose  pressure  triceps  insulin
## mean      3.8450521 120.894531 69.1054688 20.5364583 79.79948
## Desv.Estandar 3.3695781 31.972618 19.3558072 15.9522176 115.24400
## Mediana      3.0000000 117.000000 72.0000000 23.0000000 30.50000
## IQR          5.0000000 41.250000 18.0000000 32.0000000 127.25000
## Min          0.0000000 0.000000 0.0000000 0.0000000 0.00000
## Max          17.0000000 199.000000 122.0000000 99.0000000 846.00000
## Rango         17.0000000 199.000000 122.0000000 99.0000000 846.00000
## Cuartil1      1.0000000 99.000000 62.0000000 0.0000000 0.00000
## Cuartil3      6.0000000 140.250000 80.0000000 32.0000000 127.25000
## N             768.0000000 768.000000 768.0000000 768.0000000 768.00000
## ErrorEstandar 0.1215892 1.153712 0.6984425 0.5756261 4.15851
## IC95MediaLower 3.6067373 118.633255 67.7365214 19.4082312 71.64880
## IC95MediaUpper 4.0833669 123.155808 70.4744161 21.6646854 87.95016
## Varianza      11.3540563 1022.248314 374.6472712 254.4732453 13281.18008
##      mass  pedigree  age
## mean      31.9925781 0.47187630 33.2408854
## Desv.Estandar 7.8841603 0.33132860 11.7602315
## Mediana      32.0000000 0.37250000 29.0000000
## IQR          9.3000000 0.38250000 17.0000000
## Min          0.0000000 0.07800000 21.0000000
## Max          67.1000000 2.42000000 81.0000000
## Rango         67.1000000 2.34200000 60.0000000
## Cuartil1      27.3000000 0.24375000 24.0000000
## Cuartil3      36.6000000 0.62625000 41.0000000
## N             768.0000000 768.0000000 768.0000000
## ErrorEstandar 0.2844951 0.01195579 0.4243608
## IC95MediaLower 31.4349677 0.44844295 32.4091382
## IC95MediaUpper 32.5501886 0.49530965 34.0726326
## Varianza      62.1599840 0.10977864 138.3030459
```

0.2.3 2.3. Analysys Normal Distribution

a. Pregnant

```
# Normality Test "Pregnant Vairable"
```


```
df <- dataset [1]
dvector= df$pregnant
shapiro.test(dvector)
```

```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.90428, p-value < 2.2e-16
```

For a 95% confidence interval the variable does not follow a normal distribution

Graphycally:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df), color = "#FF6666", add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quartiles") +
  theme_minimal() +
  ggtitle("QQ PLOT pregnant") +
  theme(plot.title = element_text(hjust = 0.5))
```



```
df <- dataset [1]
dvector= df$pregnant
```

b. Glucose

```
# Normality Test "Glucose Vairable"
```

```
df <- dataset [2]
dvector= df$glucose
shapiro.test(dvector)
```


```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.9701, p-value = 1.986e-11
```

For a 95% confidence interval the variable does not follow a normal distribution

Graphycally:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df), color = "#FF6666", add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quartiles") +
  theme_minimal() +
  ggtitle("QQ PLOT glucose") +
  theme(plot.title = element_text(hjust = 0.5))
```

QQ PLOT glucose



```
df <- dataset [2]
dvector= df$glucose
```

c. Pressure

```
# Normality Test "Pressure Variable"
```


```
df <- dataset [3]
dvector= df$pressure
shapiro.test(dvector)
```

```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.81892, p-value < 2.2e-16
```

For a 95% confidence interval the variable does not follow a normal distribution Graphycally:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df),color = "#FF6666",add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quantiles") +
  theme_minimal() +
  ggtitle("QQ PLOT pressure") +
  theme(plot.title = element_text(hjust = 0.5))
```

QQ PLOT pressure



```
df <- dataset [3]
dvector= df$pressure
```

d. Triceps

```
# Normality Test "Triceps Variable"
```

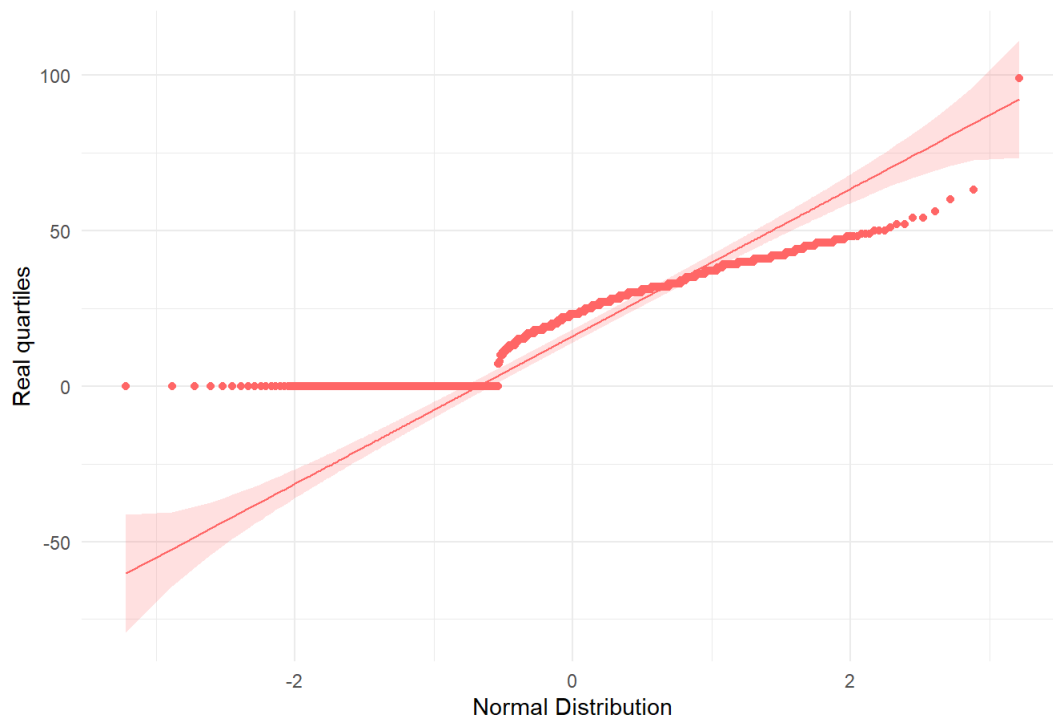
```
df <- dataset [4]
dvector= df$triceps
shapiro.test(dvector)
```

```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.90463, p-value < 2.2e-16
```

For a 95% confidence interval the variable does not follow a normal distribution Graphically:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df), color = "#FF6666", add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quartiles") +
  theme_minimal() +
  ggtitle("QQ PLOT triceps") +
  theme(plot.title = element_text(hjust = 0.5))
```

QQ PLOT triceps



```
df <- dataset [4]
dvector= df$triceps
```

e. Insulin

```
# Normality Test "Insulin Variable"
```


```
df <- dataset [5]
dvector= df$insulin
shapiro.test(dvector)
```

```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.72202, p-value < 2.2e-16
```

For a 95% confidence interval the variable does not follow a normal distribution Graphycally:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df),color = "#FF6666",add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quartiles") +
  theme_minimal() +
  ggtitle("QQ PLOT insulin") +
  theme(plot.title = element_text(hjust = 0.5))
```


QQ PLOT insulin



```
df <- dataset [5]
dvector= df$insulin
```

f. Mass

```
# Normality Test "Mass Variable"
```

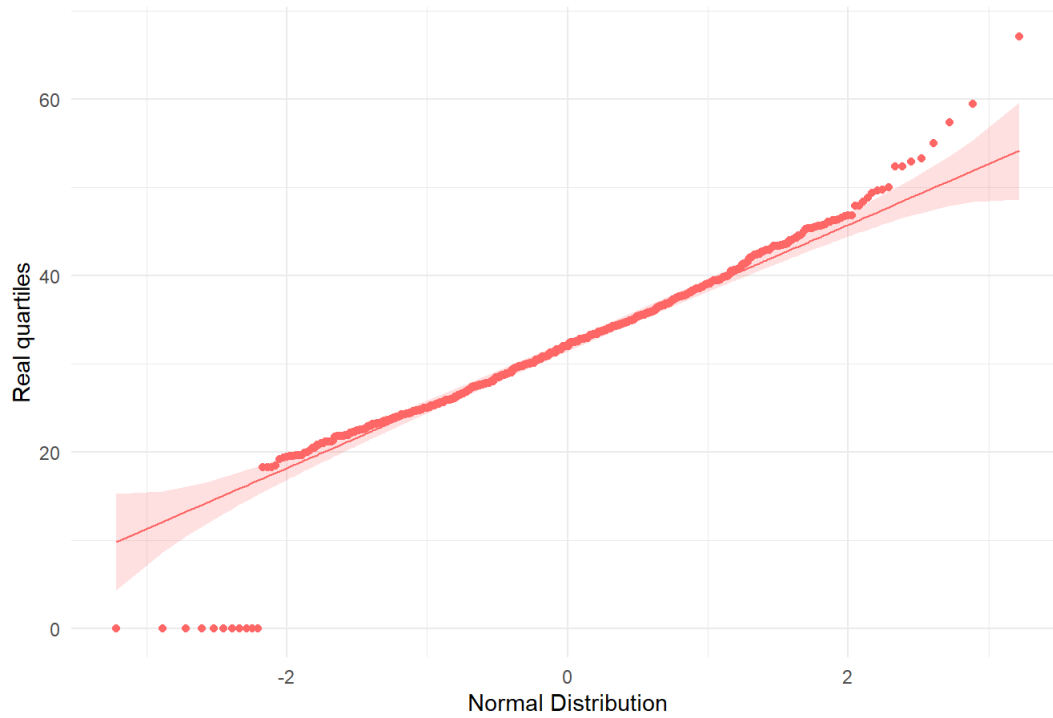
```
df <- dataset [6]
dvector= df$mass
shapiro.test(dvector)
```

```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.94999, p-value = 1.842e-15
```

For a 95% confidence interval the variable does not follow a normal distribution Graphycally:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df),color = "#FF6666",add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quartiles") +
  theme_minimal() +
  ggtitle("QQ PLOT mass") +
  theme(plot.title = element_text(hjust = 0.5))
```

QQ PLOT mass



```
df <- dataset [6]
dvector= df$mass
```

g. Pedigree

```
# Normality Test "Pedigree Variable"
```


```
df <- dataset [7]
dvector= df$pedigree
shapiro.test(dvector)
```

```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.83652, p-value < 2.2e-16
```

For a 95% confidence interval the variable does not follow a normal distribution Graphycally:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df),color = "#FF6666",add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quartiles") +
  theme_minimal() +
  ggtitle("QQ PLOT pedigree") +
  theme(plot.title = element_text(hjust = 0.5))
```

QQ PLOT pedigree



```
df <- dataset [7]
dvector= df$pedigree
```

h. Age

```
# Normality Test "Pedigree Variable"
```


```
df <- dataset [8]
dvector= df$age
shapiro.test(dvector)
```

```
##
## Shapiro-Wilk normality test
##
## data: dvector
## W = 0.87477, p-value < 2.2e-16
```

For a 95% confidence interval the variable does not follow a normal distribution Graphycally:

```
#...hacemos el qqplot
name1 <- names(df)
ggqqplot(df, x = names(df),color = "#FF6666",add.params = list(color = "black"))+
  xlab("Normal Distribution") + ylab("Real quartiles") +
  theme_minimal() +
  ggtitle("QQ PLOT age") +
  theme(plot.title = element_text(hjust = 0.5))
```

QQ PLOT age



```
df <- dataset [8]
dvector= df$age
```

0.2.4 2.4. Analysis of statistically significant differences.

As the variables do not follow a normal distribution, we will use non-parametric techniques.

a. Brown-Forsythe Test, to check if there is equality of variances

```
bf.test(pregnant ~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : pregnant and diabetes
##
## statistic : 34.89219
## num df : 1
## denom df : 455.9558
## p.value : 6.821926e-09
##
## Result : Difference is statistically significant.
## -----
```

```
bf.test(glucose ~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : glucose and diabetes
##
## statistic : 189.1048
## num df : 1
## denom df : 461.3317
## p.value : 2.644161e-36
##
## Result : Difference is statistically significant.
## -----
```

```
bf.test(pressure~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : pressure and diabetes
##
## statistic : 2.934666
## num df   : 1
## denom df : 471.3066
## p.value  : 0.08735425
##
## Result   : Difference is not statistically significant.
## -----
```

```
bf.test(triceps ~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : triceps and diabetes
##
## statistic : 3.883182
## num df   : 1
## denom df : 472.1018
## p.value  : 0.04935586
##
## Result   : Difference is statistically significant.
## -----
```

```
bf.test(insulin ~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : insulin and diabetes
##
## statistic : 10.89591
## num df   : 1
## denom df : 415.753
## p.value  : 0.001046929
##
## Result   : Difference is statistically significant.
## -----
```

```
bf.test(mass ~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : mass and diabetes
##
## statistic : 74.29262
## num df   : 1
## denom df : 573.4725
## p.value  : 6.566238e-17
##
## Result   : Difference is statistically significant.
## -----
```

```
bf.test(pedigree ~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : pedigree and diabetes
##
## statistic : 20.94721
## num df : 1
## denom df : 454.5111
## p.value : 6.100481e-06
##
## Result : Difference is statistically significant.
## -----
```

```
bf.test(age~ diabetes, data = dataset)
```

```
##
## Brown-Forsythe Test (alpha = 0.05)
## -----
## data : age and diabetes
##
## statistic : 47.89662
## num df : 1
## denom df : 575.7773
## p.value : 1.201513e-11
##
## Result : Difference is statistically significant.
## -----
```

There are difference statistically significant excepto in pression variable

b. Mann-Whitney Test to check if there are equility of means.

```
wilcox.test(pregnant~ diabetes, data = dataset, paired=FALSE)
```

```
##
## Wilcoxon rank sum test with continuity correction
##
## data: pregnant by diabetes
## W = 50985, p-value = 3.745e-08
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(glucose~ diabetes, data = dataset, paired=FALSE)
```

```
##
## Wilcoxon rank sum test with continuity correction
##
## data: glucose by diabetes
## W = 28391, p-value < 2.2e-16
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(pressure~ diabetes, data = dataset, paired=FALSE)
```

```
##
## Wilcoxon rank sum test with continuity correction
##
## data: pressure by diabetes
## W = 55415, p-value = 7.559e-05
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(triceps~ diabetes, data = dataset, paired=FALSE)
```

```
##
## Wilcoxon rank sum test with continuity correction
##
## data: triceps by diabetes
## W = 59814, p-value = 0.01296
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(insulin~ diabetes, data = dataset, paired=FALSE)
```

```
##  
## Wilcoxon rank sum test with continuity correction  
##  
## data: insulin by diabetes  
## W = 61927, p-value = 0.06566  
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(mass~ diabetes, data = dataset, paired=FALSE)
```

```
##  
## Wilcoxon rank sum test with continuity correction  
##  
## data: mass by diabetes  
## W = 41866, p-value < 2.2e-16  
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(pedigree~ diabetes, data = dataset, paired=FALSE)
```

```
##  
## Wilcoxon rank sum test with continuity correction  
##  
## data: pedigree by diabetes  
## W = 52769, p-value = 1.197e-06  
## alternative hypothesis: true location shift is not equal to 0
```

```
wilcox.test(age~ diabetes, data = dataset, paired=FALSE)
```

```
##  
## Wilcoxon rank sum test with continuity correction  
##  
## data: age by diabetes  
## W = 41950, p-value < 2.2e-16  
## alternative hypothesis: true location shift is not equal to 0
```

Only the insulin variable does not present statistically significant difference in equality of means for a 95% confidence interval.

As there are no variables that do not present statistically significant differences in both tests, we will use all the variables in the models.

0.2.5 2.5. Generate a partition of 70% of the sample for training, reserving 30% for validation.

```
# define an 70%/30train/test split of the dataset  
split=0.70  
  
trainIndex <- createDataPartition(dataset$diabetes, p=split, list=FALSE)  
data_train <- dataset[ trainIndex,]  
data_test <- dataset[-trainIndex,]
```

The training sample is made up of 538 observations and the test sample of 230 observations.

0.2.6 2.6. Cross validation will be applied to the 70% partition

```
# set-up test options  
control <- trainControl(method="cv", number=5)  
seed <- (338)  
metric <- "Accuracy"
```

0.2.7 2.7. Training algorithms

I am going to train the following algorithms: a) Linear Discriminant Analysis b) Logistic Regression c) Knn d) CART e) Bagged CART f) Random Forest g) Stochastic Gradient Boosting

```
# train algorithms
```

```
# Linear Discriminant Analysis
```

```
set.seed(seed)
fit.lda <- train(diabetes~., data=data_train, method="lda", metric=metric, preProc=c("center", "scale"), trControl=control)
```

```
# Logistic Regression
```

```
set.seed(seed)
fit.glm <- train(diabetes~., data=data_train, method="glm", metric=metric, trControl=control)
```

```
# kNN
```

```
set.seed(seed)
fit.knn <- train(diabetes~., data=data_train, method="knn", metric=metric, preProc=c("center", "scale"), trControl=control)
```

```
# CART
```

```
set.seed(seed)
fit.cart <- train(diabetes~., data=data_train, method="rpart", metric=metric, trControl=control)
```

```
# Bagged CART
```

```
set.seed(seed)
fit.treebag <- train(diabetes~., data=data_train, method="treebag", metric=metric, trControl=control)
```

```
# Random Forest
```

```
set.seed(seed)
fit.rf <- train(diabetes~., data=data_train, method="rf", metric=metric, trControl=control)
```

```
# Stochastic Gradient Boosting (Generalized Boosted Modeling)
```

```
set.seed(seed)
fit.gbm <- train(diabetes~., data=data_train, method="gbm", metric=metric, trControl=control, verbose=FALSE)
```

```
# Compare algorithms
```

```
results <- resamples(list(lda=fit.lda, logistic=fit.glm,
  knn=fit.knn, cart=fit.cart,
  bagged_CART=fit.treebag, rf=fit.rf, Stochastic_Gradient_Boosting=fit.gbm))
```

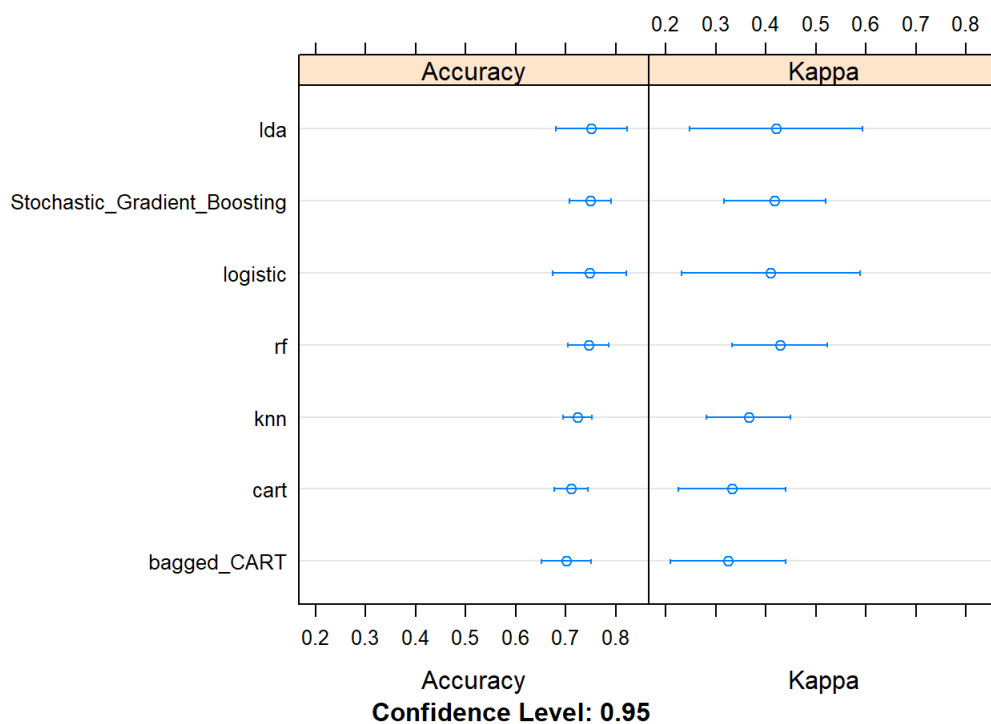
```
# Table comparison
```

```
summary(results)
```

```
##
## Call:
## summary.resamples(object = results)
##
## Models: lda, logistic, knn, cart, bagged_CART, rf, Stochastic_Gradient_Boosting
## Number of resamples: 5
##
## Accuracy
##           Min. 1st Qu.  Median    Mean 3rd Qu.
## lda          0.6666667 0.7383178 0.7407407 0.7510211 0.7962963
## logistic      0.6666667 0.7222222 0.7383178 0.7473174 0.7962963
## knn           0.6851852 0.7196262 0.7314815 0.7230876 0.7314815
## cart          0.6759259 0.6944444 0.7102804 0.7101073 0.7222222
## bagged_CART    0.6481481 0.6759259 0.7009346 0.7008307 0.7314815
## rf            0.7037037 0.7222222 0.7500000 0.7454656 0.7663551
## Stochastic_Gradient_Boosting 0.7037037 0.7383178 0.7500000 0.7490654 0.7570093
##           Max. NA's
## lda          0.8130841  0
## logistic      0.8130841  0
## knn           0.7476636  0
## cart          0.7476636  0
## bagged_CART    0.7476636  0
## rf            0.7850467  0
## Stochastic_Gradient_Boosting 0.7962963  0
##
## Kappa
##           Min. 1st Qu.  Median    Mean 3rd Qu.
## lda          0.2123177 0.3952000 0.3986351 0.4203290 0.5308057
## logistic      0.2123177 0.3435981 0.3986351 0.4100086 0.5308057
## knn           0.2560778 0.3470301 0.3930233 0.3655550 0.4077156
## cart          0.2190083 0.2732463 0.3601896 0.3325615 0.3716613
## bagged_CART    0.2095532 0.2674419 0.3127258 0.3242892 0.3930233
## rf            0.3260530 0.3759630 0.4485628 0.4279508 0.4735288
## Stochastic_Gradient_Boosting 0.3175355 0.3737458 0.4205087 0.4180083 0.4415897
##           Max. NA's
## lda          0.5646867  0
## logistic      0.5646867  0
## knn           0.4239282  0
## cart          0.4387022  0
## bagged_CART    0.4387022  0
## rf            0.5156465  0
## Stochastic_Gradient_Boosting 0.5366615  0
```

I will do a doplot to see the results of the algorithm training graphically


```
# Dot-plot comparison
dotplot(results)
```



1 3. Resultados

We make the confusion matrix for each of the algorithms on the test sample

```
# Matriz de confusion con datos de validacion
print("Accuracy Test Data Linear Discriminant Analysis")
```

```
## [1] "Accuracy Test Data Linear Discriminant Analysis"
```

```
pred <- predict(newdata=data_test,fit.Ida)
real <- as.factor(data_test$diabetes)
cm_train_tot <- caret::confusionMatrix(data=pred,reference=real)
print(cm_train_tot)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction neg pos
##      neg 132  30
##      pos  18  50
##
##      Accuracy : 0.7913
##      95% CI : (0.733, 0.8419)
##      No Information Rate : 0.6522
##      P-Value [Acc > NIR] : 2.878e-06
##
##      Kappa : 0.5233
##
##      Mcnemar's Test P-Value : 0.1124
##
##      Sensitivity : 0.8800
##      Specificity : 0.6250
##      Pos Pred Value : 0.8148
##      Neg Pred Value : 0.7353
##      Prevalence : 0.6522
##      Detection Rate : 0.5739
##      Detection Prevalence : 0.7043
##      Balanced Accuracy : 0.7525
##
##      'Positive' Class : neg
##
```

```
# Matriz de confusion con datos de validacion
print("Accuracy Test Data Linear Logistic Regression")
```

```
## [1] "Accuracy Test Data Linear Logistic Regression"
```

```
pred <- predict(newdata=data_test,fit.glm)
real <- as.factor(data_test$diabetes)
cm_train_tot <- caret::confusionMatrix(data=pred,reference=real)
print(cm_train_tot)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction neg pos
##      neg 132  29
##      pos  18  51
##
##      Accuracy : 0.7957
##      95% CI : (0.7377, 0.8458)
##      No Information Rate : 0.6522
##      P-Value [Acc > NIR] : 1.383e-06
##
##      Kappa : 0.5347
##
##      Mcnemar's Test P-Value : 0.1447
##
##      Sensitivity : 0.8800
##      Specificity : 0.6375
##      Pos Pred Value : 0.8199
##      Neg Pred Value : 0.7391
##      Prevalence : 0.6522
##      Detection Rate : 0.5739
##      Detection Prevalence : 0.7000
##      Balanced Accuracy : 0.7588
##
##      'Positive' Class : neg
##
```

```
# Matriz de confusion con datos de validacion
print("Accuracy Test Data Knn")
```

```
## [1] "Accuracy Test Data Knn"
```

```
pred <- predict(newdata=data_test,fit.knn)
real <- as.factor(data_test$diabetes)
cm_train_tot <- caret::confusionMatrix(data=pred,reference=real)
print(cm_train_tot)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction neg pos
##      neg 127 33
##      pos  23 47
##
##      Accuracy : 0.7565
##      95% CI : (0.6958, 0.8105)
##      No Information Rate : 0.6522
##      P-Value [Acc > NIR] : 0.000421
##
##      Kappa : 0.4472
##
## Mcnemar's Test P-Value : 0.229102
##
##      Sensitivity : 0.8467
##      Specificity : 0.5875
##      Pos Pred Value : 0.7938
##      Neg Pred Value : 0.6714
##      Prevalence : 0.6522
##      Detection Rate : 0.5522
##      Detection Prevalence : 0.6957
##      Balanced Accuracy : 0.7171
##
##      'Positive' Class : neg
##
```

```
# Matriz de confusion con datos de validacion
print("Accuracy Test Data Cart")
```

```
## [1] "Accuracy Test Data Cart"
```

```
pred <- predict(newdata=data_test,fit.cart)
real <- as.factor(data_test$diabetes)
cm_train_tot <- caret::confusionMatrix(data=pred,reference=real)
print(cm_train_tot)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction neg pos
##      neg 134 39
##      pos  16 41
##
##      Accuracy : 0.7609
##      95% CI : (0.7004, 0.8145)
##      No Information Rate : 0.6522
##      P-Value [Acc > NIR] : 0.000245
##
##      Kappa : 0.435
##
## Mcnemar's Test P-Value : 0.003012
##
##      Sensitivity : 0.8933
##      Specificity : 0.5125
##      Pos Pred Value : 0.7746
##      Neg Pred Value : 0.7193
##      Prevalence : 0.6522
##      Detection Rate : 0.5826
##      Detection Prevalence : 0.7522
##      Balanced Accuracy : 0.7029
##
##      'Positive' Class : neg
##
```

```
# Matriz de confusion con datos de validacion
print("Accuracy Test Data Bagged Cart")
```

```
## [1] "Accuracy Test Data Bagged Cart"
```

```

pred <- predict(newdata=data_test,fit.treebag)
real <- as.factor(data_test$diabetes)
cm_train_tot <- caret::confusionMatrix(data=pred,reference=real)
print(cm_train_tot)

```

```

## Confusion Matrix and Statistics
##
##      Reference
## Prediction neg pos
##      neg 126  26
##      pos  24  54
##
##      Accuracy : 0.7826
##      95% CI : (0.7236, 0.8341)
##      No Information Rate : 0.6522
##      P-Value [Acc > NIR] : 1.156e-05
##
##      Kappa : 0.518
##
##      Mcnemar's Test P-Value : 0.8875
##
##      Sensitivity : 0.8400
##      Specificity : 0.6750
##      Pos Pred Value : 0.8289
##      Neg Pred Value : 0.6923
##      Prevalence : 0.6522
##      Detection Rate : 0.5478
##      Detection Prevalence : 0.6609
##      Balanced Accuracy : 0.7575
##
##      'Positive' Class : neg
##

```

```

# Matriz de confusion con datos de validacion
print("Accuracy Test Data Random Forest")

```

```

## [1] "Accuracy Test Data Random Forest"

```

```

pred <- predict(newdata=data_test,fit.rf)
real <- as.factor(data_test$diabetes)
cm_train_tot <- caret::confusionMatrix(data=pred,reference=real)
print(cm_train_tot)

```

```

## Confusion Matrix and Statistics
##
##      Reference
## Prediction neg pos
##      neg 128  25
##      pos  22  55
##
##      Accuracy : 0.7957
##      95% CI : (0.7377, 0.8458)
##      No Information Rate : 0.6522
##      P-Value [Acc > NIR] : 1.383e-06
##
##      Kappa : 0.5456
##
##      Mcnemar's Test P-Value : 0.7705
##
##      Sensitivity : 0.8533
##      Specificity : 0.6875
##      Pos Pred Value : 0.8366
##      Neg Pred Value : 0.7143
##      Prevalence : 0.6522
##      Detection Rate : 0.5565
##      Detection Prevalence : 0.6652
##      Balanced Accuracy : 0.7704
##
##      'Positive' Class : neg
##

```

```
# Matriz de confusion con datos de validacion
print("Accuracy Test Data Stochastic Gradient Boosting")
```

```
## [1] "Accuracy Test Data Stochastic Gradient Boosting"
```

```
pred <- predict(newdata=data_test,fit.gbm)
real <- as.factor(data_test$diabetes)
cm_train_tot <- caret::confusionMatrix(data=pred,reference=real)
print(cm_train_tot)
```

```
## Confusion Matrix and Statistics
##
##      Reference
## Prediction neg pos
##      neg 130  28
##      pos  20  52
##
##      Accuracy : 0.7913
##      95% CI : (0.733, 0.8419)
##      No Information Rate : 0.6522
##      P-Value [Acc > NIR] : 2.878e-06
##
##      Kappa : 0.529
##
##      Mcnemar's Test P-Value : 0.3123
##
##      Sensitivity : 0.8667
##      Specificity : 0.6500
##      Pos Pred Value : 0.8228
##      Neg Pred Value : 0.7222
##      Prevalence : 0.6522
##      Detection Rate : 0.5652
##      Detection Prevalence : 0.6870
##      Balanced Accuracy : 0.7583
##
##      'Positive' Class : neg
##
```

1.1 4.- Final Conclusions and limitations

As final conclusions, it is worth highlighting the good performance of all the models in the TRUE POSITIVE RATE (in this case the detection of negative cases of diabetes), exceeding in all cases the 80% success rate, which gives great security in the negative prediction of diabetes.

While the specificity in some cases is not as high as would be desired. This may be due to the fact that the dataset is unbalanced, with negative cases that exceed 3 times the rate of positive cases.

For future lines of research, the datasets could be balanced using artificial sampling techniques in order to balance the predictions. Another possibility would be with the use of the ROC curve, to estimate a different cut-off point.